

KEVAL SAVALIYA

Agentic AI | RAG Systems | AI Automation
skeval1601@gmail.com | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)
Stuttgart, Germany

PROFESSIONAL SUMMARY

AI Systems Engineer with 2+ years building production RAG pipelines, multi-agent architectures, and LLM-powered automation workflows. MSc student in Automotive Software Engineering (TU Chemnitz) with hands-on experience deploying agentic retrieval systems, MCP-based tool integrations, and hybrid vector search infrastructure. Seeking a Master's thesis in RAG, Agentic AI, or AI Automation to apply production-grade engineering experience to applied research problems.

EDUCATION

Master of Science in Automotive Software Engineering April 2024 – Present
Technische Universität Chemnitz, Germany GPA: 2.16/5.0

- Coursework: Neurocomputing, Image Understanding, AUTOSAR Architecture, Embedded Systems, Software Service Engineering

Bachelor of Science in Information Technology June 2020 – May 2023
Gujarat University, India GPA: 3.0/5.0

- Coursework: Advanced Python Programming, Data Structures for Analytics, Statistical Methods, Database Systems, Machine Learning Fundamentals

TECHNICAL SKILLS

RAG & Retrieval	Hybrid Search (Dense + Sparse), BM25, pgvector, ChromaDB, Pinecone, FAISS, Cohere Reranker, BGE-M3, SentenceTransformers
Agentic AI	LangGraph, LangChain Agents, MCP Servers, Multi-Agent Architectures, Pydantic, Tool Use
LLMs & Eval	Claude API, OpenAI, Gemini, Groq, Ragas, DeepEval, LangSmith, Prompt Engineering
Automation	n8n Workflows, Apify, Apollo.io, Airtable, Lead Generation Pipelines
Programming	Python, SQL, C/C++, Java
ML & Data	PyTorch, TensorFlow, scikit-learn, XGBoost, Pandas, NumPy, BERT, Transformers
Development	FastAPI, Flask, Streamlit, REST APIs, Docker, GitHub Actions, CI/CD
Cloud & Infra	AWS EC2, GCP, RabbitMQ, PostgreSQL, Kubernetes, Jenkins
Languages	English (C1), German (B1)

PROFESSIONAL EXPERIENCE

IT Innovation & Digitalization Intern September 2026 – Present
MAHLE International GmbH Stuttgart, Germany

- Working within the IT Innovation & Digitalization department on AI-driven process automation and intelligent workflow solutions
- Applying RAG pipeline and LLM integration experience to enterprise-scale digitalization challenges in the automotive sector

AI Automation Engineer (Part-time) August 2026 – Present
Soccerkinetics Stuttgart, Germany (Remote)

- Built automated video processing pipelines combining computer vision (object detection) and local LLM inference for content analysis and classification
- Designed and deployed end-to-end automation workflows reducing manual content production effort by over 80%

AI Engineer Intern April 2026 – August 2026
Mindioo Prague, Czech Republic (Remote)

- Designed and deployed a production NCF recommendation pipeline on AWS EC2 using Docker, RabbitMQ multi-queue messaging, and PostgreSQL

- Built a RAG QA agent using LangChain, ChromaDB, and TinyLlama to answer knowledge-base queries with grounded, cited responses
- Developed LLM-based features and prompt engineering pipelines in collaboration with the CTO

Data Scientist Intern

July 2025 – September 2025

Data Glacier

Remote, India

- Built Python automation scripts and KPI dashboards, reducing manual reporting time by 60% and model evaluation processing time by 30%
- Improved ML model accuracy by 15% through advanced feature engineering and SQL-based data quality audits maintaining 98% pipeline consistency

Data Scientist Intern

June 2023 – January 2024

HackSec Infotech

India

- Contributed to CI/CD pipeline deployment using Jenkins, Docker, and Kubernetes, reducing deployment time by 30%
- Analysed large-scale datasets using Python and SQL to identify trends; maintained 95% data quality standards through rigorous preprocessing workflows

KEY PROJECTS

AutoDoc Intelligence – Multi-Agent RAG System for Automotive Documents

LangGraph, LlamaIndex, BGE-M3, pgvector, BM25, Cohere Reranker, Ragas, DeepEval, FastAPI, Docker, AWS EC2

[GitHub](#)

- Built a multi-agent RAG system over automotive technical documents (ISO 9001, IATF 16949) with hybrid dense + sparse retrieval (BGE-M3 + BM25 + Cohere reranker), improving context precision from 0.61 to 0.81 on a 100-query golden eval set
- Designed a LangGraph router dispatching queries across specialised agents (QA, compliance conflict detection, escalation) with LangSmith tracing and p95 latency under 1.5s
- Implemented a Ragas + DeepEval CI gate in GitHub Actions blocking PRs above 5% hallucination rate, with eval results auto-posted as PR comments

CreoBot – AI-Powered SaaS Content Platform

FastAPI, Next.js 14, Groq LLaMA 3.3 70B, pgvector, Supabase, Stripe

[Live Demo](#)

- Shipped a production SaaS application with Stripe subscription billing, pgvector semantic search, and Groq-powered LLM content generation serving real users
- Architected a FastAPI backend with Supabase (PostgreSQL + Auth) and a Next.js 14 frontend deployed on Vercel with full CI/CD

MCP AI Developer Tools – Clinical & Workflow Automation

Python, TypeScript, MCP SDK, FastMCP, Pydantic

[GitHub](#)

- Built two production MCP servers exposing 7 AI tools (code review, commit generation, IEC 62304 compliance, FHIR resource generation) integrable with Claude and any MCP-compatible client
- Combined both servers into a domain-specific Claude Skill; built a TypeScript MCP client connecting to the Python server for cross-language agentic integration

CERTIFICATIONS

Generative AI Mastermind – Outskill **Generative AI: Prompt Engineering Basics – IBM via Coursera**

ADDITIONAL INFORMATION

Work Authorization: Student visa holder in Germany, authorized to work part-time during studies

Availability: Available for Master's Thesis supervision in RAG, Agentic AI, or AI Automation